



Audio Engineering Society

Convention Paper 7963

Presented at the 128th Convention
2010 May 22–25 London, UK

The papers at this Convention have been selected on the basis of a submitted abstract and extended precis that have been peer reviewed by at least two qualified anonymous reviewers. This convention paper has been reproduced from the author's advance manuscript, without editing, corrections, or consideration by the Review Board. The AES takes no responsibility for the contents. Additional papers may be obtained by sending request and remittance to Audio Engineering Society, 60 East 42nd Street, New York, New York 10165-2520, USA; also see www.aes.org. All rights reserved. Reproduction of this paper, or any portion thereof, is not permitted without direct permission from the Journal of the Audio Engineering Society.

A GPGPU Approach to Improved Acoustic Finite Difference Time Domain Calculations

Jamie A S Angus¹, Andrew Counce²

School of Computing, Science and Engineering, University of Salford, Salford, Greater Manchester,
M5 4WT, United Kingdom

¹j.a.s.angus@salford.ac.uk

²a.j.counce@student.salford.ac.uk

ABSTRACT

This paper shows how to improve the efficiency and accuracy of Finite Difference Time Domain acoustic simulation by both calculating the differences using spectral methods and performing these calculations on a Graphics Processing Unit (GPU) rather than a CPU. These changes to the calculation method result in an increase in accuracy as well as a reduction in computational expense. The recent advances in the way that GPU's are programmed (for example using CUDA on Nvidia's GPU) now make them an ideal platform on which to perform scientific computations at very high speeds and very low power consumption.

1. INTRODUCTION

Finite difference time domain methods were first conceived for electrodynamics problems in the 1960s by Yee [1] in order to solve Maxwell's equations for electrodynamics problems. His work was built upon by and the term FDTD coined by Alan Taflove in 1980 [2]. Meloney and Cummings first applied this method to acoustics in 1995, [3].

The finite difference time domain method involves discretising both time and space, and then solving the 2 dimensional wave equation by way of the leapfrog method. In order for this to work one must separate

acoustic particle velocity (in both the x and y directions as particle velocity is a vector quantity) and pressure (which is scalar) onto separate points on the grid as shown in Figure 1.




Figure 1 FDTD Grid used in simulations (Adapted from picture in [4])

1.1. Spectral Methods

Spectral Methods form one member of the three major technologies used in the solution of partial differential equations. These are the finite element method, the aforementioned finite difference methods, and spectral methods.

If one is solving a PDE where the initial data is smooth enough, then a spectral method offers very high accuracy or lower computation times [5]. Spectral methods for solving differential equations involve using the Fourier transform to perform the differentiation. This is achieved by taking the impulse response of a perfect differentiator and convolving it with the input data. This convolution can be done by multiplying the two 'signals' in the 'frequency' domain (or k domain in the case of the input FDTD data). This means that the whole extent of the input signal is taken into consideration when calculating the derivative as opposed to just a few surrounding points, thus removing the dispersion error created by the leapfrog method of solving the PDE's. This has the advantage that the grid would no longer have to be sampled at higher than the normal Nyquist limit and therefore could be sampled at just two times the highest frequency, reducing 2D calculations by up to 5 times and 3D calculations by up to 25 times.

1.2. GPU Computing

The architecture of the GPU differs greatly from that of the CPU in that the CPU will have up to 4 separate processing cores, whereas the GPU has up to (as of 2009) 200+. This means that if the problem can be naturally split into numerous operations that can be calculated in parallel (i.e. Each calculation can be calculated independently of the other calculations being carried out) then intuitively substantial speedups can be achieved.

This lends itself to the world of scientific computing where many problems can be effectively parallelized, a good example being the fast Fourier transform. This has achieved major speedups over the equivalent CPU algorithms [6]. This indicates a huge increase in computational efficiency meaning that higher degrees of accuracy can be achieved in more feasible amounts of time.

The major limitations of GPU programming are the time taken to access global memory and the size of the local memory and cache. This means that in order to achieve maximum efficiency accesses to the global or host memory should be kept to an absolute minimum (ideally eliminated altogether) this means that it is advantageous to recalculate values as they are needed as opposed to calculating once and then accessing the memory. In other words in order to maximise the efficiency of the program one should optimise the arithmetic intensity to minimise the memory accesses. Thus using the spectral methods means that although there are potentially more calculation steps ($N \log_2 N$) more to be exact, there should be a lot less memory accesses due to the much smaller grid size than is necessary for the straight leap frog formula.

1.3. CUDA

CUDA, Complete Unified Digital Architecture, is NVIDIA's implementation of GPGPU. CUDA is based around the industry standard C programming language and provides a simple way of programming for GPUs. The programming interface remains fairly high level with the compiler and runtime deciding which thread to run on which processor etc., meaning that the programmer is left with a very intuitive way to program [7].

2. STANDARD FDTD CALCULATION METHOD ON GPU

The Finite Difference Time Difference method, when applied to acoustics, is in essence a finite solution to the second order wave equation below, and Euler's equation to relate pressure to velocity for simplicities sake this will be explained in terms of the 1 dimensional case.

$$\text{1D Wave Equation: } \frac{\partial^2 p}{\partial x^2} - \frac{1}{c^2} \frac{\partial^2 p}{\partial t^2} = 0 \quad (1)$$

$$\text{Differentiated Euler's Formula: } \rho_0 \frac{\partial}{\partial x} \left[\frac{\partial u}{\partial t} \right] = \frac{\partial^2 p}{\partial t^2} \quad (2)$$

The above equation (1) is then equated to Euler's formula, which has been differentiated with respect to space (2).

$$\rho_0 \frac{\partial}{\partial x} \left[\frac{\partial u}{\partial t} \right] = \frac{1}{c^2} \frac{\partial^2 p}{\partial t^2} \quad (3)$$

Integrating equation (3) with respect to time yields:

$$\frac{\partial p}{\partial t} = -\rho_0 c \frac{\partial u}{\partial x} \quad (4)$$

Expressing this equation and (2) in terms of both time and space then discretising both pressure and velocity onto separate points as shown in the Figure 1 yields:

$$p\left(x_{i+\frac{1}{2}}, t_{n+\frac{1}{2}}\right) = p\left(x_{i+\frac{1}{2}}, t_{n-\frac{1}{2}}\right) - \rho_0 c \frac{\Delta t}{\Delta x} [u(x_{i+1}, t_n) - u(x_i, t_n)] \quad (5)$$

$$u(x_i, t_n) = p(x_i, t_{n-1}) - \frac{1}{\rho_0} \frac{\Delta t}{\Delta x} \left[u\left(x_{i+\frac{1}{2}}, t_{n-\frac{1}{2}}\right) - u\left(x_{i-\frac{1}{2}}, t_{n-\frac{1}{2}}\right) \right] \quad (6)$$

These equations are now dependent only on the previous values pressure and velocity respectively.

2.1. Sources

There are three types of source that are useful in calculating FDTD when applied directly to a pressure node, these are:

1. Hard Source
2. Soft Source
3. Transparent Source

2.1.1. Hard Source

This is implemented by, at the desired point of excitation, the FDTD equation is negated and the pressure of that point (x_{sx}) is excited only by the required source equation.

$$p\left(x_{sx}, t_{n+\frac{1}{2}}\right) = S(x_{sx}) \quad (7)$$

This type of source is useful for the fact that the signal that inputs into the grid is identical to that desired. However because the FDTD equations have been negated, the source will cause reflections and scattering.

2.1.2. Soft Source

A soft source is implemented by adding the source data onto the calculated FDTD pressure in the form:

$$p\left(x_{sx}, t_{n+\frac{1}{2}}\right) = S(x_{sx}) + p\left(x_{i+\frac{1}{2}}, t_{n-\frac{1}{2}}\right) - \rho_0 c \frac{\Delta t}{\Delta x} [u(x_{i+1}, t_n) - u(x_i, t_n)] \quad (8)$$

This type of source has the advantage of not introducing reflections or scattering however the downfall is that the signal is changed to a degree by the grid. The reason for this is that the grid acts as a form of medium and as such has an impulse response that alters the signal being fed into the grid.

2.1.3. Transparent Source

The transparent source aims to reduce the warping caused by the impedance of the grid, by matching the impedance of the source to that of the grid. This is achieved by convolving the source signal with the

impulse response of the grid. In two dimensions this is impossible due to the high frequency limit of the grid (Section 3.3.2); this means that a proper impulse response cannot be achieved, as a Dirac delta function will not give good enough data.

Due to the complexity of this method and the fact that due to the aforementioned limitations the effect of this source is to only reduce the warping, not eliminate it, this method is impractical for this proof of concept style project.

2.2. Boundaries

Boundaries in the calculation of FDTD are subject to a large amount of research [8] and can be of many levels of complexity. The only two to be used in this project are a perfectly hard boundary, and a perfectly matched layer.

2.2.1. Perfectly Hard Boundary

This is achieved by way of; at the specific node where a boundary is required the velocity is forced to be zero. This simulates a surface with a reflection coefficient of one and an absorption coefficient of zero. This is useful when trying to model the effects of diffraction and other low frequency effects although of limited use in practical simulations of real spaces.

2.2.2. Perfectly Matched Layer

A perfectly matched layer is a necessary piece of technology in any FDTD calculation, in that it is used to absorb any energy heading for the edge of the grid. It works to gradually absorb the energy of a wave propagating in any direction.




Figure 2 Demonstration of Perfectly Matched Layer

This means that it can be used to model perfectly anechoic conditions and also, when a boundary that is not totally absorbent is used, some energy is transmitted through the boundary, a PML can be used

to absorb this and prevent spurious reflections from entering the grid.

The PML is implemented by way of an auxiliary parameter, which is defined for a PML at both edges of a 1D grid as;

let d = PML Depth

$$\alpha = \begin{cases} i < d & \alpha = \frac{1}{3} \left(\frac{i}{d} \right)^3 \\ d < i < N - d & \alpha = 0 \\ i > N - d & \alpha = \frac{1}{3} \left(\frac{N - i}{d} \right)^3 \end{cases} \quad i = 1, 2, 3, \dots, N - 1 \quad (9)$$

Then the FDTD equations are in order to accommodate this auxiliary parameter thus:

$$p \left(x_{i+\frac{1}{2}}, t_{n+\frac{1}{2}} \right) = p \left(x_{i+\frac{1}{2}}, t_{n-\frac{1}{2}} \right) \left(\frac{1 - \alpha_i}{1 + \alpha_i} \right) - \rho_0 c \frac{\Delta t}{\Delta x} \left(\frac{1}{1 + \alpha_i} \right) [u(x_{i+1}, t_n) - u(x_i, t_n)] \quad (10)$$

$$u(x_i, t_n) = p(x_i, t_{n-1}) \left(\frac{1 - \alpha_i}{1 + \alpha_i} \right) - \frac{1}{\rho_0} \frac{\Delta t}{\Delta x} \left(\frac{1}{1 + \alpha_i} \right) [u(x_{i+\frac{1}{2}}, t_{n-\frac{1}{2}}) - u(x_{i-\frac{1}{2}}, t_{n-\frac{1}{2}})] \quad (11)$$

2.3. Adaptations for Implementation in CUDA

The first major change is that instead of looping through the values in space, an individual thread is launched to calculate each point. For each time step there are 2 sets of threads launched, this is to ensure that all of the velocity values are calculated before the pressures begin to be calculated. This means that the pressure and velocity values of each point at each time step are calculated at the same time, greatly reducing the time taken iterating through the points in space.

For example in calculating the perfectly matched layer, the values of the coefficients are calculated in each thread, thus saving an allocation and transfer of three large 2D, $N \times N$, arrays into the device memory. In order to reduce memory transfers between the host and the device only the output pressure after each time step will be copied back in to the host memory and output. However if velocity were required this

would have to be copied as well as or in place of pressure.

2.4. Proof of Accurate Calculation

The following output clearly shows the FDTD working and showing diffraction working, the source was the first half of a sine wave (similar to a cosine window), to represent an impulse. The subsequent plots show that the wave is propagating across the grid and becoming incident upon a wall with holes in it.




Figure 3 Plots of pressure versus position showing the propagation of the wave over time

Also shown in this plot is the PML, it is clearly visible that there are no spurious reflections from the boundaries as the wave is absorbed entirely by the PML, thus accurately modeling a perfectly reflection free environment.

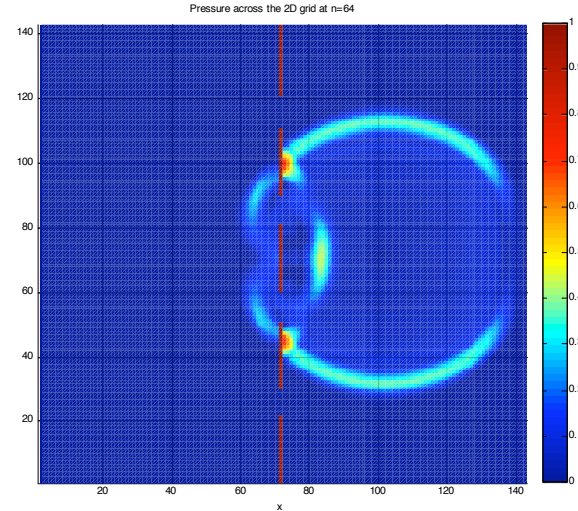


Figure 4 Output of FDTD program showing diffusion taking place.

It is clear from the above output that the FDTD is accurately modeling the diffraction of the wave around the edges of the barriers. The red lines indicated a perfectly hard surface (velocity forced to 0). A line has been drawn between the excitation point and the edge of one of the hard surfaces. This is to demonstrate that diffraction is taking place. Any energy within the white lines, i.e. in the shadow zone (shaded), can only be there due to diffraction, in this example we see plenty of energy in this region proving that diffraction is being modeled effectively.

2.5. Problems with this Method

The major problem with the standard method of calculating the FDTD is that there is a high frequency limit to the method. This is demonstrated in the figure 5. In order to accurately calculate the frequency response the impulse response of a low pass filter ($\text{sinc}(\omega_0 t)$) is used with the cut off frequency set to where the desired highest frequency is.




Figure 5 Demonstration of upper frequency limit of FDTD Grid

The above figures show that the grid begins to seriously distort the grid at frequencies approaching $0.2f_s$, this means that in order to model a range specific of frequencies, with a reasonable degree of accuracy, we would have to sample the grid at $5f_s$ in the case of the audible range this equates to sampling at $>200kHz$, this leads to a huge amount of data and calculation due to the dependency between the spatial and time discretisations, otherwise known as the Courant limit [9]. This states that for a square grid (i.e. $dx=dy$)

$$\Delta t \leq \frac{\Delta x}{c\sqrt{2}} \quad (12)$$

This shows the direct proportionality between Δx and Δt , thus the smaller Δt becomes the proportionally smaller Δx must become as well, thus in addition to more time steps being necessary as ($\Delta t = 1/f_s$) there are also more spatial points required.

3. IMPLEMENTATION OF SPECTRAL METHODS IN CUDA

3.1. Theory

The most basic way of thinking about the spectral method used in this case is that rather than calculating the derivative of the velocity or pressure from the adjoining values one can replace this with the appropriate value from a differentiation matrix created from the spectral differentiation. This has the advantage of rather than evaluating just two points in order to calculate the derivative at that point, all of the points in the data set are being evaluated producing the theoretically most accurate derivative

possible with the N points. In one dimension the frequency domain representation of the differentiation matrix takes the form:

$$w(i) = \begin{cases} i < \frac{N}{2} & y = i \\ i = \frac{N}{2} & y = 0 \\ i > \frac{N}{2} & y = i - N \end{cases} \quad i = 0, 1, 2, \dots, N-1 \quad (13)$$

This equation applied to when $N=30$ is shown in figure 6.




Figure 6 Differentiation Matrix for 1D spectral differentiation.

This impulse response of the perfect differentiator is then multiplied by the result of the Fourier Transform of the input 'signal' to form the differential of the input. This can then be used instead of the centre differences in the FDTD equations (5) and (6) in order to achieve a more accurate estimation of the state of the FDTD grid at that point:

$$p\left(x_{i+\frac{1}{2}}, t_{n+\frac{1}{2}}\right) = p\left(x_{i+\frac{1}{2}}, t_{n-\frac{1}{2}}\right) \left(\frac{1-\alpha_i}{1+\alpha_i}\right) - \rho_0 c \frac{\Delta t}{\Delta x} \left(\frac{1}{1+\alpha_i}\right) w_u(x_i) \quad (14)$$

$$u(x_i, t_n) = p(x_i, t_{n-1}) \left(\frac{1-\alpha_i}{1+\alpha_i}\right) - \frac{1}{\rho_0} \frac{\Delta t}{\Delta x} \left(\frac{1}{1+\alpha_i}\right) w_p(x_{i+\frac{1}{2}}) \quad (15)$$

In order to validate the differentiation method, the derivative of a sine wave was calculated using this method, the results are shown in figure 7 plotted against the perfect derivative $\cos(x)$:




Figure 7 Graph showing the calculated spectral derivative and the theoretical perfect derivative.

This calculation was performed in MATLAB as opposed to in CUDA, therefore it has been performed with double precision floating point numbers as opposed to if the calculation was performed in CUDA it would have been performed in single precision floating point numbers. The error shown on the above plot is somewhere close to the limit of the accuracy of the floating-point algebra. This is phenomenally accurate differentiation more accurate than could have been achieved by way of centre differences. However the downfall with spectral methods is the accuracy of the result depends hugely on the smoothness of the function. For example see the spectral differentiation of a triangular window shown in figure 8.




Figure 8 Graphs showing a triangular window and its spectral derivative (Based on Illustration in [5])

3.2. Considerations in CUDA

The implementation in CUDA is very similar in general to that for the more conventional method of performing FDTD calculations. In that there will be a **for** loop to iterate through the time steps then a set of CUDA kernels for both pressure and velocity.

The major difference is the addition of the FFTs. Rather than the author writing a bespoke FFT, the NVIDIA CUFFT library was used. This is a very general purpose FFT algorithm and as such is potentially not the most efficient FFT that could be used for this purpose, however the CUFFT library has been shown by benchmark tests to be more efficient for large transform sizes than the ubiquitous FFTW [6]. The reason that the CUFFT has no benefit over small transform sizes is simply due to the fact that CPUs have larger registers and therefore can store the necessary parts of the problem more locally to the core than in the somewhat limited registers of a GPU core.

The arrays used in this type of calculation are potentially very large, especially when the possibility of 3D simulations is considered, therefore it was considered prudent to use the CUFFT library and

accept the potential lack of speed on small scale simulations.

Waterfall Plot of a 1D FDTD Calculated by Spectral Method




Figure 9 Waterfall plot of 1D FDTD calculated with Spectral Methods

3.3. Results

Figure 9 shows a lovely wave propagating across the grid with minimal ripples. This suggests that a 2D spectral FDTD would be more accurate and have fewer errors than the standard centre difference method of calculating FDTD.

4. CONCLUSIONS

4.1. Finite Difference Time Domain on GPGPU

Presented is a proof of concept that Finite Difference Time Domain calculations can be carried out with ease on a Graphics Processing Unit thanks to the advent of CUDA. These calculations, although not always as accurate perhaps as their equivalent on a CPU due to the lack of double precision support (although this point is irrelevant for the latest GPUs that do support double precision), the calculations are

carried out very well and the power of the method is in no way compromised.

4.2. Spectral Methods

Spectral methods are a phenomenally powerful way of solving partial differential equations. The results obtained in this, albeit, limited investigation are very impressive and the potential for this method is impressive. The power of this method definitely warrants further work.

5. FURTHER WORK

5.1. Spectral Methods

The further work required is to expand the spectral method into 2 or even 3 dimensions and to test that the high frequency limit is, as theorised, less restrictive than for the more general way to calculate FDTD.

5.2. GPGPU

The limited time available for this project meant that the decision was taken to take no time optimising the CUDA code and testing the speed at which the code executes. This is defiantly a project for further investigation as the code used to execute here was laid out in the most intuitive way as opposed to the potentially fastest and most efficient way.

6. REFERENCES

- [1] K. Yee, "Numerical solutions of initial boundary value problems involving Maxwell's equations in isotropic media," IEEE Transactions on Antennas and Propagation., vol. AP-14, pp. 302–307, 1966.
- [2] A. Taflove, "Application of the finite-difference time-domain method to sinusoidal steady-state electromagnetic penetration problems," IEEE Transactions on Electromagnetic Compatibility, vol. 22, pp. 191-202, Aug. 1980.
- [3] J.G. Meloney and K.E. Cummings, "Adaption of FDTD techniques to acoustic modelling", 11th Annual Review Prog Applied Computational Electromagnetics, Vol.2, 724 (1995)
- [4] Ostashev, V. E. & Wilson, D. K. & Liu, L. & Aldridge D. F. & Symons, N. P. & Marlin, D.: 2005 *.Equations for Finite-Difference, Time-Domain Simulation of Sound Propagation in Moving inhomogeneous Media and Numerical Implementation.*, J. Acoust. Soc. Am. **117** (2).
- [5] L. N. Trefethen, *Spectral Methods in MATLAB*, SIAM, Philadelphia, 2000.
- [6] N. K. Govindaraju, et al. (2008). 'High performance discrete Fourier transforms on graphics processors'. In SC '08: Proceedings of the 2008 ACM/IEEE conference on Supercomputing, pp. 1-12, Piscataway, NJ, USA. IEEE Press.
- [7] NVIDIA Corporation, "CUDA Programming Guide" NVIDIA, Santa Clara CA, 2007, text available at www.nvidia.com/cuda.
- [8] Hyok Jeong, Ian Drumm, Brian Horner, Yiu Wai Lam, "The modeling frequency dependent boundary conditions in FDTD simulation of concert hall acoustics", Proceedings of the 19th ICA, Madrid, Spain. 2007
- [9] A. Taflove and S. C. Hagness, *Computational Electrodynamics: The Finite-Difference Time-Domain Method*, 2nd ed. Norwood, MA: Artech House, 2000.

7. APPENDICIES

7.1. 2D FDTD

```
1 //Include files
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <cuda_runtime.h>
5 #include <cutil.h>
6 #include <math.h>
7
8 //Initialise CUDA
9 #if __DEVICE_EMULATION__
10
11 bool InitCUDA(void){return true;}
12
13 #else
14 bool InitCUDA(void)
15 {
16     int count = 0;
17     int i = 0;
18
19     cudaGetDeviceCount(&count);
20     if(count == 0) {
21         fprintf(stderr, "There is no device.\n");
22         return false;
23     }
24
25     for(i = 0; i < count; i++) {
26         cudaDeviceProp prop;
27         if(cudaGetDeviceProperties(&prop, i) == cudaSuccess) {
28             if(prop.major >= 1) {
29                 break;
30             }
31         }
32     }
33     if(i == count) {
```

```
34         fprintf(stderr, "There is no device supporting CUDA.\n");
35         return false;
36     }
37     cudaSetDevice(i);
38
39     printf("CUDA initialized.\n");
40     return true;
41 }
42
43 #endif
44
45
46 //Declare pressure and velocity arrays
47
48
49 const int N = 143;
50
51
52
53
54
55 __global__ void CalculateVelocities(int n, float* pDevice, size_t pitch, float* uxDev, float* uyDev, int
56 xSize, int ySize, int PMLDepth)
57 {
58     //Declare Constants
59     float rho0=1.21, c=343, fs=20000;
60     float dt=1/fs;
61     float dx=dt*c*sqrt(2.0);
62     float pConst = rho0*pow(c, 2)*(dt/dx);
63     float uConst = (1/rho0)*(dt/dx);
64     //Set up indexes
65     int x = (blockIdx.x*blockDim.x+threadIdx.x)+1;
66     int y = (blockIdx.y*blockDim.y+threadIdx.y)+1;
67     float alphaUx2, alphaUy2, alphaP2;
68
69
70     if (x<PMLDepth)
```

```
71         alphaUx2=(1/3.0)*pow(((float) (PMLDepth-x-1)/PMLDepth),3);
72     else if (x>xSize+PMLDepth)
73         alphaUx2=(1/3.0)*pow((float) (((float)x-(xSize+PMLDepth))/PMLDepth),3);
74     else
75         alphaUx2=0;
76
77     //Calculate y velocity values for Alpha
78     if (y<PMLDepth)
79         alphaUy2=(1/3.0)*pow((float) ((float) (PMLDepth-y-1)/PMLDepth),3);
80     else if (y>ySize+PMLDepth)
81         alphaUy2=(1/3.0)*pow((float) (((float)y-(ySize+PMLDepth))/PMLDepth),3);
82     else
83         alphaUy2=0;
84
85     alphaP2 = (alphaUx2+alphaUy2)/2;
86
87     if (n==1){
88         uxDev[y*pitch/4+x]=0;
89         uyDev[y*pitch/4+x]=0;}
90
91     if (y==2)
92     y=y;
93     if ((y<N-2&& x<N) && (y>1&& x>1)){
94         //Calculate Velocities
95         uxDev[y*pitch/4+x]=uxDev[y*pitch/4+x]*((1.0-alphaUx2)/(1.0+alphaUx2))-
96         uConst*(1.0/(1.0+alphaP2))*(pDevice[y*pitch/4+x]-pDevice[y*pitch/4+(x-1)]);
97         uyDev[y*pitch/4+x]=uyDev[y*pitch/4+x]*((1.0-alphaUy2)/(1.0+alphaUy2))-
98         uConst*(1.0/(1.0+alphaP2))*(pDevice[y*pitch/4+x]-pDevice[(y-1)*pitch/4+x]);
99
100         ///barrier to diffract round
101         if ((x<40&&y==40)|| (x<40&&y==60)|| (x==30&&y<60&&y>40)){
102             uxDev[y*pitch/4+x]=0;
103             uyDev[y*pitch/4+x]=0;
104         }
105
106
107     }
```

```
108 }
109
110 __global__ void CalculatePressures(int n, float* pDevice, size_t pitch, float* uxDev, float* uyDev, int
111 xSize, int ySize, int PMLDepth )
112 {
113
114     float rho0=1.21, c=343, fs=20000;
115     float dt=1/fs;
116     float dx=dt*c*sqrt(2.0);
117     float pConst = rho0*pow(c,2)*(dt/dx);
118     float uConst = (1/rho0)*(dt/dx);
119     //Set up indexes
120     int x = blockIdx.x*blockDim.x+threadIdx.x;
121     int y = blockIdx.y*blockDim.y+threadIdx.y;
122
123     float alphaUx2, alphaUy2, alphaP2;
124
125     if (x<PMLDepth)
126         alphaUx2=(1/3.0)*pow(((float)(PMLDepth-x-1)/PMLDepth),3);
127     else if (x>xSize+PMLDepth)
128         alphaUx2=(1/3.0)*pow((float)((float)x-(xSize+PMLDepth))/PMLDepth),3);
129     else
130         alphaUx2=0;
131
132     //Calculate y velocity values for Alpha
133     if (y<PMLDepth)
134         alphaUy2=(1/3.0)*pow((float)((float)(PMLDepth-y-1)/PMLDepth),3);
135     else if (y>ySize+PMLDepth)
136         alphaUy2=(1/3.0)*pow((float)((float)y-(ySize+PMLDepth))/PMLDepth),3);
137     else
138         alphaUy2=0;
139
140     alphaP2 = (alphaUx2+alphaUy2)/2;
141
142     if (n==1)
143         pDevice[y*pitch/4+x]=0;
144
```

```

145
146     if ((y<N-2&& x<N) && (y>1&& x>1))
147     {
148         //Calculate Pressure Output
149         ////Apply Excitation
150         pDevice[y*pitch/4+x]=pDevice[y*pitch/4+x]*((1.0-alphaP2)/(1.0+alphaP2))
151             - (pConst*(1.0/(1.0+alphaUx2))* (uxDev[y*pitch/4+(x+1)] -
152 uxDev[y*pitch/4+x])
153             +pConst*(1.0/(1.0+alphaUy2))* (uyDev[(y+1)*pitch/4+x] -
154 uyDev[y*pitch/4+x]));
155         if (x==71&& y==50)
156         {
157             if ((2.0*3.142*900*n*dt)<3.142)
158                 pDevice[y*pitch/4+x]=3*sin(2.0*3.142*900*n*dt); /// (2.0*3.142*900*n*dt);
159             else
160                 pDevice[y*pitch/4+x]=0;
161         }
162     }
163
164
165
166 }
167
168 extern "C" int Run2DFDTD(float pOutput2[N][N],int timeSteps)
169 {
170     //Declare Constants
171     const int PMLDepth = 21;
172     const float rho0=1.21,c=343,fs=20000;
173     const float dt=1/fs;
174     const float dx=dt*c*sqrt(2.0);
175     //float pOutput2[N][N];
176
177     const int H=3,W=3;
178     const int xSize = 101, ySize=101;
179
180     float alphaP[N][N],alphaUx[N][N],alphaUy[N][N];           //alpha matrices
181

```



```
182     float *pDev,*uxDev,*uyDev;
183     size_t pitch;
184
185     //Calculate pressure alpha
186
187
188
189     //
190     //alphaP[1][1]=(alphaUx[1][1]+alphaUy[1][1])/2;
191
192     #pragma region Memory Allocations
193     size_t arraySize = N*sizeof(float);
194     cudaMallocPitch((void**)&pDev, &pitch,arraySize, arraySize);
195     cudaMallocPitch((void**)&uxDev, &pitch,arraySize, arraySize);
196     cudaMallocPitch((void**)&uyDev, &pitch,arraySize, arraySize);
197 #pragma endregion
198
199
200     dim3 dimBlock(1,1);
201     dim3 dimGrid(N/dimBlock.y, N / dimBlock.x);
202     //loop through time steps
203     int n = 0;
204     for (n=1;n<timeSteps;n++)
205     {
206         //Invoke kernel to calculate the grid
207         CalculateVelocities<<<dimGrid,dimBlock>>>(n,pDev,pitch,uxDev,uyDev,xSize,ySize,PMLDepth);
208         CalculatePressures<<<dimGrid,dimBlock>>>(n,pDev,pitch,uxDev,uyDev,xSize,ySize,PMLDepth);
209
210         float devMemOutput[N][N];
211
212         printf("n=%i Completed!!!\n",n);
213         if (n==timeSteps-1)
214         {
215             cudaMemcpy2D (devMemOutput, N*sizeof(float), pDev, pitch,
216                           N*sizeof(float), N*sizeof(float),
217                           cudaMemcpyDeviceToHost);
218             printf("Results Copy Complete!!");
```

```
219         for (int i = 0;i<N;i++)
220         {
221             for (int j=0;j<N;j++)
222             {
223                 pOutput2[i][j]=devMemOutput[i][j];
224             }
225         }
226     }
227 }
228
229
230 cudaFree (pDev) ;
231 cudaFree (uxDev) ;
232 cudaFree (uyDev) ;
233
234
235 return (1) ;
236 }
```

7.2. Spectral FDTD

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <cuda_runtime.h>
4  #include <cutil.h>
5  #include < cufft.h>
6  #include <math.h>
7
8
9  const cufftComplex j={0,1};
10 const float pi = 3.142;
11 #define CUFFT_FORWARD -1
12 #define CUFFT_INVERSE 1
13 const float c = 343.0;
14 const float fs = 10000.0;
15 /*****
16  /* Init CUDA */
17  *****/
18 #if __DEVICE_EMULATION__
19
20 bool InitCUDA(void){return true;}
21
22 #else
23 bool InitCUDA(void)
24 {
25     int count = 0;
26     int i = 0;
27
28     cudaGetDeviceCount(&count);
29     if(count == 0) {
30         fprintf(stderr, "There is no device.\n");
31         return false;
32     }
33
34     for(i = 0; i < count; i++) {
35         cudaDeviceProp prop;
```

```
36         if(cudaGetDeviceProperties(&prop, i) == cudaSuccess) {
37             if(prop.major >= 1) {
38                 break;
39             }
40         }
41     }
42     if(i == count) {
43         fprintf(stderr, "There is no device supporting CUDA.\n");
44         return false;
45     }
46     cudaSetDevice(i);
47
48     printf("CUDA initialized.\n");
49     return true;
50 }
51
52 #endif
53 /*****
54  /* Kernels
55  *****/
56 __device__ cufftComplex CMul(cufftComplex z,cufftComplex x)
57 {
58     cufftComplex answer;
59     answer.x=z.x*x.x-z.y*x.y;
60     answer.y=(z.x+z.y)*(x.x+x.y)-z.x*x.x-z.y*x.y;
61     return(answer);
62 }
63 __global__ void SetArray(cufftComplex *a,bool isP,int N)
64 {
65     int k = threadIdx.x;
66     float h=2*3.142/(float)N;
67
68     //if(isP)
69     //    a[k].x=3*(float)exp(-5*pow((h*k-1)-1,2));
70     //    //a[k]=fmax(0,1-fabs(h*k-3.142)/2);
71     //else
72         a[k].x=0;
```

```
73
74
75
76 }
77 __global__ void ZeroComplex(cufftComplex *a)
78 {
79     int k = threadIdx.x;
80     a[k].x=0;
81     a[k].y=0;
82 }
83
84 __global__ void CalculateWHat(float *diffMat,cufftComplex *aHat, cufftComplex *wHat)
85 {
86     int k = threadIdx.x+1;
87     cufftComplex j = {0,1};
88     //cast diffmat as a complex number for the multiplication
89     cufftComplex temp = {diffMat[k],0};
90     //calculate wHat
91     wHat[k]=CMul(CMul(j,aHat[k]),temp);
92     k=k;
93 }
94
95 __global__ void CalcDiffMat(float *diffMatrix,int N)
96 {
97     //Calculate and write the impulse response of a perfect differentiation matrix of size 1*N
98     int k = threadIdx.x;
99     if(k<N/2)
100         diffMatrix[k]=k;
101     if(k==N/2)
102         diffMatrix[k]=0;
103     if(k>N/2)
104         diffMatrix[k]=k-N;
105
106
107
108 }
```

```

109 __global__ void CalcP(cufftComplex *p,cufftComplex *uDif,float pConst,int N,int i,int timeSteps,int
110 PMLDepth)
111 {
112     int k = threadIdx.x;
113     if(k==fabs(N/2.0))
114         k=k;
115     float alpha;
116     int gridSize=N-2*PMLDepth;
117
118     if (k<PMLDepth)
119         alpha=(1/3.0)*pow(((float) (PMLDepth-k)/PMLDepth),3);
120     else if (k>gridSize+PMLDepth)
121         alpha=(1/3.0)*pow(((float) (k-(gridSize+PMLDepth))/PMLDepth),3);
122     else
123         alpha=0;
124
125
126     p[k].x=p[k].x*((1.0-alpha)/(1.0+alpha))-pConst*(1.0/(1.0+alpha))*(uDif[k].x/(3.142*N));
127
128
129     float pulseWidth=10,pulseCentreX=40,pulseCentreT=5;
130     if(i==1&& k>(N/2-pulseWidth/2)&&k<(N/2+pulseWidth/2))
131         //p[k].x=1;
132         p[k].x=p[k].x+(1-(1.93*cos(2*pi*((k-(N/2-pulseWidth/2))/pulseWidth))
133             +1.29*cos(4*pi*((k-(N/2-pulseWidth/2))/pulseWidth))
134             -0.388*cos(6*pi*((k-(N/2-pulseWidth/2))/pulseWidth))
135             +0.0322*cos(8*pi*((k-(N/2-pulseWidth/2))/pulseWidth))));
136
137 }
138 __global__ void CalcU(cufftComplex *u,cufftComplex *pDif,float uConst,int N,int PMLDepth)
139 {
140     int k = threadIdx.x+1;
141
142     float alpha;
143     int gridSize=N-2*PMLDepth;
144     //calculate alpha value for this point
145     if (k<PMLDepth)

```



```
146         alpha=(1/3.0)*pow(((float) (PMLDepth-k)/PMLDepth),3);
147     else if (k>gridSize+PMLDepth)
148         alpha=(1/3.0)*pow(((float) ((float)k-(gridSize+PMLDepth))/PMLDepth),3);
149     else
150         alpha=0;
151
152
153     u[k].x=u[k].x*((1.0-alpha)/(1.0+alpha))-uConst*(1.0/(1.0+alpha))*(pDif[k].x/(3.142*N));
154
155
156 }
157
158 /*****
159 /* Main Functions                                     */
160 /*****
161 int main(int argc, char* argv[])
162 {
163
164     if(!InitCUDA()) {
165         return 0;
166     }
167     unsigned int timer = 0;
168     CUT_SAFE_CALL( cutCreateTimer( &timer));
169     CUT_SAFE_CALL( cutStartTimer( timer));
170
171     //Declare Physical Constants
172     float rho0=1.21;
173
174     float dt=1/(2*fs);          //-|conditions required to run
175     float dx=2*dt*c;           //-|simulation at the courant limit
176
177     float pConst = rho0*pow(c,2)*(dt/dx)*dt*c;
178     float uConst = (1/rho0)*(dt/dx)*dt*c;
179
180     //Setup Grid
181     int PMLDepth = 30;
182     float gridWidth = 2,endTime=0.03;
```

```

183     const int N = fabs(gridWidth/dx)+2*PMLDepth;
184     int timeSteps = fabs(endTime/dt);
185
186 #pragma region declare and allocate device variables
187     float *diffMat;
188     cufftComplex *uHat_d, *pHat_d, *p_d, *u_d, *uDifHat_d, *pDifHat_d, *uDif_d, *pDif_d;
189     cudaMalloc((void**) &p_d, N*sizeof(cufftComplex));
190     cudaMalloc((void**) &u_d, N*sizeof(cufftComplex));
191     cudaMalloc((void**) &diffMat, N*sizeof(cufftReal));
192     cudaMalloc((void**) &pHat_d, N*sizeof(cufftComplex));
193     cudaMalloc((void**) &uHat_d, N*sizeof(cufftComplex));
194     cudaMalloc((void**) &pDif_d, N*sizeof(cufftComplex));
195     cudaMalloc((void**) &uDif_d, N*sizeof(cufftComplex));
196     cudaMalloc((void**) &pDifHat_d, N*sizeof(cufftComplex));
197     cudaMalloc((void**) &uDifHat_d, N*sizeof(cufftComplex));
198     cufftComplex* p_h=(cufftComplex*)malloc(N*sizeof(cufftComplex));
199 #pragma endregion
200
201 #pragma region Set values of input variables
202     dim3 dimBlock(1,1);
203     SetArray<<<dimBlock,N>>>(p_d,true,N);
204     ZeroComplex<<<dimBlock,N>>>(pHat_d);
205     ZeroComplex<<<dimBlock,N>>>(uHat_d);
206     ZeroComplex<<<dimBlock,N>>>(pDifHat_d);
207     ZeroComplex<<<dimBlock,N>>>(uDifHat_d);
208     ZeroComplex<<<dimBlock,N>>>(u_d);
209 #pragma endregion
210
211     //Set up differentiation matrix
212     CalcDiffMat<<<dimBlock,N>>>(diffMat,N);
213
214     FILE* input;
215     input = fopen("input.txt","w+");
216     for (int n=0;n<N;n++)
217         fprintf(input,"%f ",p_d[n]);
218     fclose(input);
219

```

```
220 #pragma region Set cufft plans
221     cufftHandle planR2C,planC2R,planC2C;
222     cufftPlan1d(&planR2C,N,CUFFT_R2C,1);
223     cufftPlan1d(&planC2R,N,CUFFT_C2R,1);
224     cufftPlan1d(&planC2C,N,CUFFT_C2C,1);
225 #pragma endregion
226
227
228
229 //Open file to record output
230 FILE* file;
231 file=fopen("output.txt","w+");
232 //loop through time steps
233
234 for (int i=0;i<timeSteps;i++)
235 {
236     //Take spectral Derivatives
237     cufftExecC2C(planC2C, p_d, pHat_d,CUFFT_FORWARD);
238     CalculateWhat<<<dimBlock,N-2>>>(diffMat,pHat_d,pDifHat_d);
239     cufftExecC2C(planC2C, pDifHat_d, pDif_d,CUFFT_INVERSE);
240
241
242 // Calculate new values of p & u
243 CalcU<<<dimBlock,N>>>(u_d,pDif_d,uConst,N,PMLDepth);
244
245     cufftExecC2C(planC2C, u_d, uHat_d,CUFFT_FORWARD);
246     CalculateWhat<<<dimBlock,N-2>>>(diffMat,uHat_d,uDifHat_d);
247     cufftExecC2C(planC2C, uDifHat_d, uDif_d,CUFFT_INVERSE);
248
249
250     CalcP<<<dimBlock,N-2>>>(p_d,uDif_d,pConst,N,i,timeSteps,PMLDepth);
251
252
253     //Write out p to output
254     cudaMemcpy(p_h,p_d,N*sizeof(cufftComplex),cudaMemcpyDeviceToHost);
255
256     //if(fmod(i,4.0)==0){
```

```
257         for (int n=0;n<N;n++)
258         {
259             fprintf(file,"%f  ",p_h[n].x);
260         }
261         fprintf(file,"\n");
262         //}
263         printf("time step n=%i Complete =D!!!\n",i);
264     }
265     fclose(file);
266     fclose(input);
267
268
269
270
271
272     CUDA_SAFE_CALL( cudaThreadSynchronize() );
273     CUT_SAFE_CALL( cutStopTimer( timer));
274     printf("Processing time: %f (ms)\n", cutGetTimerValue( timer));
275     CUT_SAFE_CALL( cutDeleteTimer( timer));
276
277     free(p_h);
278     cudaFree(p_d);
279     cudaFree(u_d);
280     cudaFree(diffMat);
281     cudaFree(pHat_d);
282     cudaFree(uHat_d);
283     cudaFree(pDif_d);
284     cudaFree(uDif_d);
285     cufftDestroy(planC2R);
286     cufftDestroy(planR2C);
287     CUT_EXIT(argc, argv);
288
289     return 0;
290 }
291
```